

Sonal Shukla (CS23176)

Lab-4 Assignment-1 : To implement a Decision Tree regression model on the Developer Performance dataset and analyze how different tree depths and splitting criteria influence model performance.

```
import pandas as pd
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
```

```
from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

```
path= "/content/drive/MyDrive/ML_DATASETS/Developer_Performance_dataset_1000 - Developer_Performance_dataset_1000.csv"
df=pd.read_csv(path)
df.head(5)
```

	Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress_Level
0	7	416	9	7	6	5.9	92	7	
1	4	269	16	13	5	5.1	85	2	
2	11	439	3	0	2	6.2	38	2	
3	8	472	15	9	4	4.2	26	5	
4	5	265	19	16	5	8.1	82	6	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df.describe()
```

	Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress_Level
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	5.840000	356.234000	9.876000	7.153000	2.961000	6.465800	56.933000	3.368000	
std	3.15854	188.15535	5.796052	5.468226	2.021278	1.439529	21.767506	2.342234	
min	1.00000	26.00000	0.000000	0.000000	0.000000	4.000000	20.000000	0.000000	
25%	3.00000	209.50000	5.000000	2.000000	1.000000	5.200000	38.000000	1.000000	
50%	6.00000	332.00000	10.000000	7.000000	3.000000	6.400000	57.000000	3.000000	
75%	9.00000	480.50000	15.000000	12.000000	5.000000	7.700000	76.000000	5.000000	
max	11.00000	993.00000	19.000000	19.000000	6.000000	9.000000	94.000000	7.000000	1

```
df.shape
```

(1000, 13)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Hours_Coding    1000 non-null   int64  
 1   Lines_of_Code   1000 non-null   int64  
 2   Bugs_Found      1000 non-null   int64  
 3   Bugs_Fixed      1000 non-null   int64  
 4   AI_Usage_Hours  1000 non-null   int64  
 5   Sleep_Hours     1000 non-null   float64 
 6   Cognitive_Load  1000 non-null   int64  
 7   Coffee_Intake   1000 non-null   int64  
 8   Stress_Level    1000 non-null   int64  
 9   Task_Duration_Hours  1000 non-null   float64 
 10  Commits         1000 non-null   int64  
 11  Errors          1000 non-null   int64  
 12  Task_Success_Rate  1000 non-null   int64  
dtypes: float64(2), int64(11)
memory usage: 101.7 KB
```

```
df.columns
```

```
Index(['Hours_Coding', 'Lines_of_Code', 'Bugs_Found', 'Bugs_Fixed',
       'AI_Usage_Hours', 'Sleep_Hours', 'Cognitive_Load', 'Coffee_Intake',
       'Stress_Level', 'Task_Duration_Hours', 'Commits', 'Errors',
       'Task_Success_Rate'],
      dtype='object')
```

```
df.isnull()
```

	Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress_Level
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
995	False	False	False	False	False	False	False	False	False
996	False	False	False	False	False	False	False	False	False
997	False	False	False	False	False	False	False	False	False
998	False	False	False	False	False	False	False	False	False
999	False	False	False	False	False	False	False	False	False

1000 rows × 13 columns

```
df.isnull().any()
```

	Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress_Level	Task_Duration_Hours	Commits	Errors	Task_Success_Rate
0	False	False	False	False	False	False	False	False	False	False	False	False	False

```
dtype: bool
```

```
df.isnull().any().any()
```

```
np.False_
```

```
df.isnull().sum()
```

```
0
Hours_Coding 0
Lines_of_Code 0
Bugs_Found 0
Bugs_Fixed 0
AI_Usage_Hours 0
Sleep_Hours 0
Cognitive_Load 0
Coffee_Intake 0
Stress_Level 0
Task_Duration_Hours 0
Commits 0
Errors 0
Task_Success_Rate 0
```

dtype: int64

```
df[df.isnull().any(axis=1)]
```

```
Hours_Coding  Lines_of_Code  Bugs_Found  Bugs_Fixed  AI_Usage_Hours  Sleep_Hours  Cognitive_Load  Coffee_Intake  Stress_Level
```

```
df.columns
```

```
Index(['Hours_Coding', 'Lines_of_Code', 'Bugs_Found', 'Bugs_Fixed',
       'AI_Usage_Hours', 'Sleep_Hours', 'Cognitive_Load', 'Coffee_Intake',
       'Stress_Level', 'Task_Duration_Hours', 'Commits', 'Errors',
       'Task_Success_Rate'],
      dtype='object')
```

```
df.shape
```

```
(1000, 13)
```

```
before_rows = df.shape[0]
df=df.dropna()
after_rows = df.shape[0]

print("Before rows: ",before_rows)
print("After Rows : ",after_rows)
print("The number of droped rows are : ", before_rows-after_rows)
```

```
Before rows: 1000
After Rows : 1000
The number of droped rows are : 0
```

```
df.shape
```

```
(1000, 13)
```

```
clean_data = df.copy()
```

```
clean_data.columns
```

```
Index(['Hours_Coding', 'Lines_of_Code', 'Bugs_Found', 'Bugs_Fixed',
       'AI_Usage_Hours', 'Sleep_Hours', 'Cognitive_Load', 'Coffee_Intake',
       'Stress_Level', 'Task_Duration_Hours', 'Commits', 'Errors',
       'Task_Success_Rate'],
      dtype='object')
```

```
clean_data.columns
```

```
Index(['Hours_Coding', 'Lines_of_Code', 'Bugs_Found', 'Bugs_Fixed',
       'AI_Usage_Hours', 'Sleep_Hours', 'Cognitive_Load', 'Coffee_Intake',
       'Stress_Level', 'Task_Duration_Hours', 'Commits', 'Errors',
       'Task_Success_Rate'],
      dtype='object')
```

```
clean_data.shape
```

```
(1000, 13)
```

```
X=clean_data.iloc[:, :-1]
print(X)
```

	Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	\
0	7	416	9	7	6	
1	4	269	16	13	5	
2	11	439	3	0	2	
3	8	472	15	9	4	
4	5	265	19	16	5	
..	
995	10	660	14	13	0	
996	9	484	13	11	1	
997	2	128	4	0	3	
998	8	266	5	3	1	
999	8	604	16	15	3	
	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress_Level	\	
0	5.9	92	7	99		
1	5.1	85	2	100		
2	6.2	38	2	55		
3	4.2	26	5	30		
4	8.1	82	6	82		
..		
995	5.5	21	2	30		
996	8.8	34	0	39		
997	8.7	42	3	49		
998	5.7	48	0	63		
999	6.9	26	4	30		
	Task_Duration_Hours	Commits	Errors			
0	10.5	20	3			
1	9.5	17	8			
2	18.3	35	2			
3	12.6	28	4			
4	7.0	25	9			
..			
995	5.2	20	5			
996	4.7	37	1			
997	2.8	5	7			
998	11.8	9	0			
999	4.4	11	1			

```
[1000 rows x 12 columns]
```

```
y=clean_data.iloc[:, -1]
print(y)
```

0	34
1	36
2	79
3	94
4	33
..	..
995	80
996	78
997	74
998	51
999	90

```
Name: Task_Success_Rate, Length: 1000, dtype: int64
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,test_size=0.3,random_state=42)
```

```
X_train
```

Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress
541	3	178	0	0	1	8.1	28	0
440	6	226	5	2	2	5.2	21	2
482	5	388	4	1	6	5.4	68	1
422	1	176	8	3	3	6.4	84	3
778	5	340	17	15	0	6.3	24	4
...
106	8	236	17	15	3	6.0	61	4
270	8	228	5	3	2	8.8	23	5
860	1	102	19	14	2	5.9	51	3
435	7	189	17	12	0	8.1	26	0
102	10	282	4	0	2	8.4	76	0

700 rows × 12 columns

Next steps: [Generate code with X_train](#) [New interactive sheet](#)**X_test**

Hours_Coding	Lines_of_Code	Bugs_Found	Bugs_Fixed	AI_Usage_Hours	Sleep_Hours	Cognitive_Load	Coffee_Intake	Stress
521	3	234	5	3	2	8.4	63	5
737	2	254	5	1	6	6.6	86	4
740	5	266	13	9	4	6.4	84	3
660	5	140	3	0	0	8.1	60	3
411	8	810	2	0	6	8.0	43	5
...
468	4	168	8	6	3	4.4	60	5
935	4	210	2	0	2	7.9	34	0
428	9	646	8	2	2	8.3	94	7
7	3	234	5	1	6	5.8	78	2
155	7	224	19	15	1	8.0	54	2

300 rows × 12 columns

Next steps: [Generate code with X_test](#) [New interactive sheet](#)**y_train**

Task_Success_Rate
541
440
482
422
778
...
106
270
860
435
102

700 rows × 1 columns

dtype: int64

y_test

Task_Success_Rate

521	39
737	36
740	36
660	49
411	82
...	...
468	49
935	77
428	30
7	40
155	49

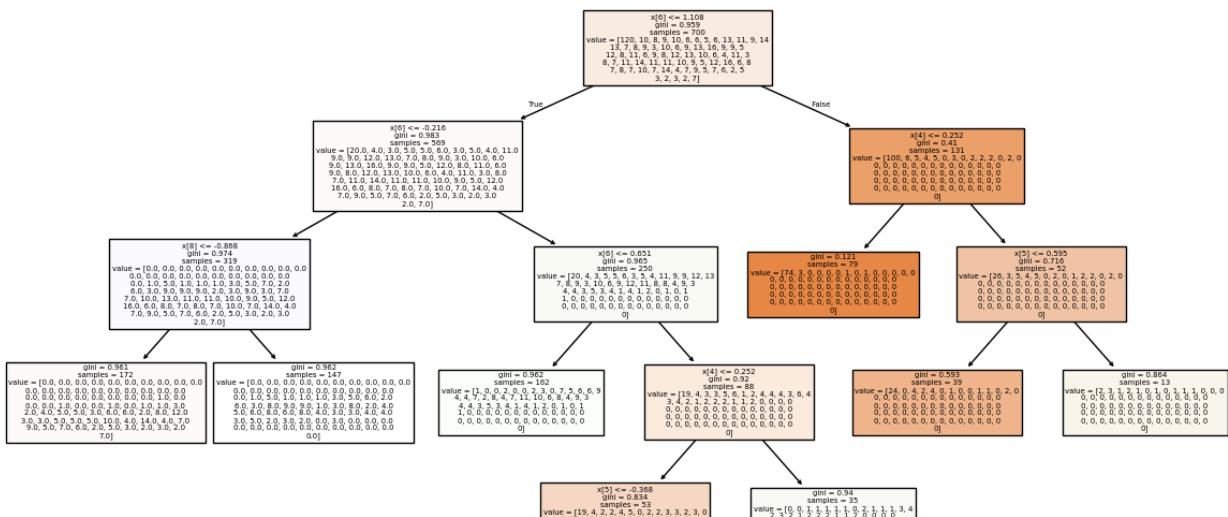
300 rows × 1 columns

dtype: int64

```
from sklearn.preprocessing import StandardScaler
scaling = StandardScaler()
X_train = scaling.fit_transform(X_train)
X_test = scaling.transform(X_test)
```

```
humidity_classifier = DecisionTreeClassifier(max_leaf_nodes= 10,random_state=0)
humidity_classifier.fit(X_train,y_train)
```

```
import matplotlib.pyplot as plt
from sklearn import tree
plt.figure(figsize=(15,10))
tree.plot_tree(humidity_classifier,filled=True)
```



```
y_pred = humidity_classifier.predict(X_test)
print(y_pred)
[50 30 51 63 30 51 63 51 30 51 63 87 51 30 51 87 87 63 43 51 51
 87 51 87 43 63 30 30 87 51 87 87 63 87 87 51 51 30 87 87 51 87 63 30 51
 51 63 51 51 87 51 51 63 63 30 51 30 63 63 87 30 87 51 30 51 87 87
 30 30 30 87 51 30 35 51 51 87 43 51 87 43 63 63 51 63 87 51 30 30 63 51
 30 30 51 51 87 63 63 43 51 63 87 87 30 30 35 51 63 51 87 30 30 63 51 51
 30 30 51 43 30 87 30 51 63 30 43 87 87 63 30 51 87 87 51 87 63 87 30
 30 63 30 31 87 51 63 30 51 51 87 51 51 87 35 63 51 51 63 51 31 63 30 35
 63 87 51 30 30 51 43 63 51 30 87 51 87 63 51 30 63 87 51 63 63 30 51
 51 30 63 51 30 51 87 30 87 51 63 63 30 87 63 63 30 63 87 63 30 51 51
 31 30 63 87 87 63 63 30 63 30 87 43 51 87 30 30 30 51 30 30 30 30 31
 87 30 87 63 51 31 51 30 87 51 51 30 30 51 63 30 43 63 87 51 63 43 51 30
 63 30 51 30 87 63 87 51 51 30 63 30 30 51 51 43 63 63 31 43 63 30
 87 30 63 87 30 51 63 51 87 30 43 51]
```

```
from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
```

```
print(confusion_matrix(y_pred,y_test))
```

```
[[51  2  1 ...  0  0  0]
 [ 2  0  1 ...  0  0  0]
 [ 0  0  0 ...  0  0  0]
 ...
 [ 0  0  0 ...  0  0  0]
 [ 0  0  0 ...  0  0  0]
 [ 0  0  0 ...  0  0  0]]
```

```
accuracy = accuracy_score(y_pred,y_test)
print(accuracy)
```

```
0.18666666666666668
```

```
print(classification_report(y_pred,y_test))
```

	precision	recall	f1-score	support
30	0.89	0.68	0.77	75
31	0.00	0.00	0.00	6
32	0.00	0.00	0.00	0
33	0.00	0.00	0.00	0